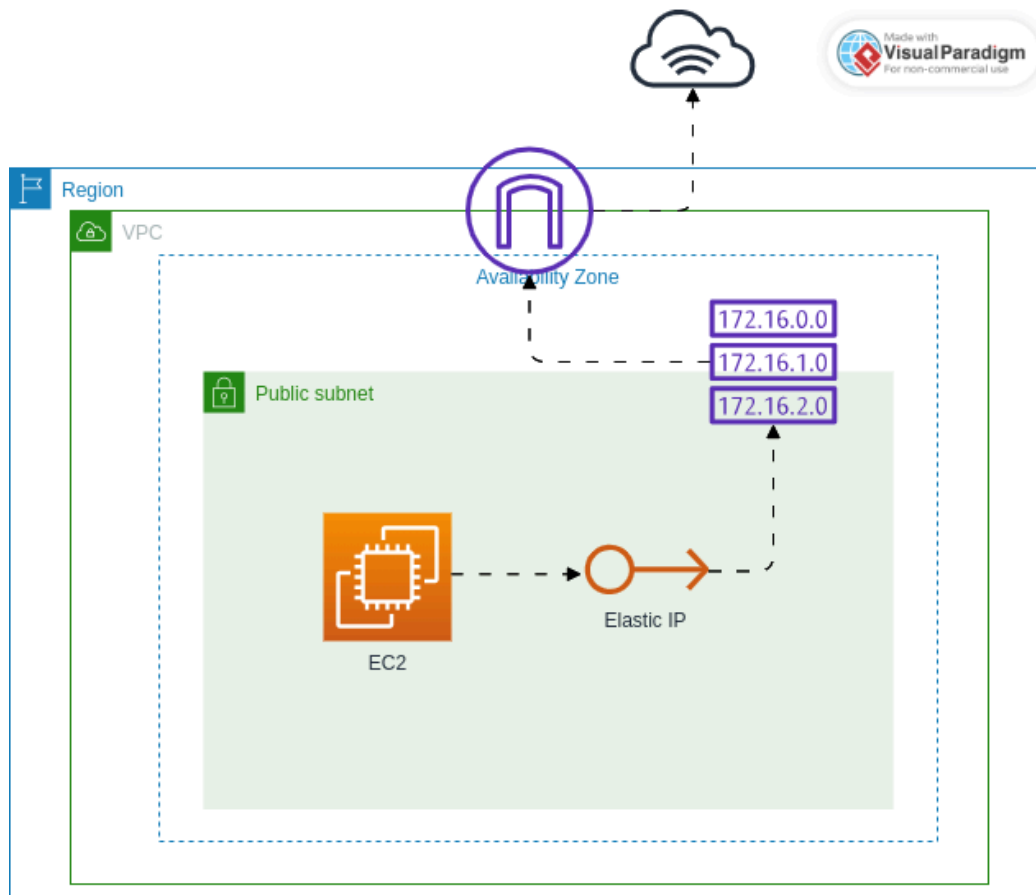


# Interview Tasks

## Task1 Cloud Infrastructure & Deployment on AWS

### Architecture Diagram



### Steps to Deploy the Application

1. Create an AWS Account
2. Create an Cloudformation template

```
# structure from here
```

```
#
```

<https://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/template-formats.html#template-comments>

**AWSTemplateFormatVersion:** 2010-09-09

**Description:** Interview Test file

**Parameters:**

**EC2InstanceSizeInput:**

**Description:** The supported instances sizes for EC2

**Type:** String

**Default:** t2.micro

**AllowedValues:**

- t3.micro
- t2.micro

**Resources:**

# VPC , subnet , igw , route table , router table to igw rule, security group , security group rules , eip, nic , ec2 , user data , ssm role , ssm policy, role assumption

#VPC

**TestVPC:**

**Type:** AWS::EC2::VPC

### Properties:

CidrBlock: 10.0.0.0/16

### Tags:

- Key: ProjectNumber

Value: 4

- Key: ProjectName

Value: interviewData

### TestIGW:

Type: AWS::EC2::InternetGateway

### Properties:

### Tags:

- Key: ProjectNumber

Value: 4

- Key: ProjectName

Value: interviewData

### TestAttachGateway:

Type: AWS::EC2::VPCEGatewayAttachment

### Properties:

VpcId: !Ref TestVPC

InternetGatewayId: !Ref TestIGW

#Subnet

TestPublicSubnet:

Type: AWS::EC2::Subnet

Properties:

VpcId: !Ref TestVPC

AvailabilityZone: "ap-south-1a"

CidrBlock: 10.0.0.1/24

Tags:

- Key: ProjectNumber

Value: 4

- Key: ProjectName

Value: interviewData

TestRouteTable:

Type: AWS::EC2::RouteTable

Properties:

VpcId: !Ref TestVPC

Tags:

- Key: ProjectNumber

Value: 4

- Key: ProjectName

Value: interviewData

TestInternetPublicRoute:

Type: AWS::EC2::Route

DependsOn: TestIGW

Properties:

RouteTableId: !Ref TestRouteTable

DestinationCidrBlock: 0.0.0.0/0

GatewayId: !Ref TestIGW

TestRouteTableToTestSubnetAssociation:

Type: AWS::EC2::SubnetRouteTableAssociation

Properties:

RouteTableId: !Ref TestRouteTable

SubnetId: !Ref TestPublicSubnet

TestInstanceSecurityGroup:

Type: AWS::EC2::SecurityGroup

## Properties:

GroupDescription: Allow EC2 traffic

VpcId: !Ref TestVPC

## SecurityGroupIngress:

- Description: Allow ssh

IpProtocol: tcp

FromPort: 22

ToPort: 22

CidrIp: 0.0.0.0/0

- Description: Allow http

IpProtocol: tcp

FromPort: 80

ToPort: 80

CidrIp: 0.0.0.0/0

- Description: Allow https

IpProtocol: tcp

FromPort: 443

ToPort: 443

CidrIp: 0.0.0.0/0

- Description: Allow all

IpProtocol: -1

CidrIp: 0.0.0.0/0

SecurityGroupEgress:

- Description: Allow ssh

IpProtocol: tcp

FromPort: 22

ToPort: 22

CidrIp: 0.0.0.0/0

- Description: Allow http

IpProtocol: tcp

FromPort: 80

ToPort: 80

CidrIp: 0.0.0.0/0

- Description: Allow https

IpProtocol: tcp

FromPort: 443

ToPort: 443

CidrIp: 0.0.0.0/0

- Description: Allow all

IpProtocol: -1

CidrIp: 0.0.0.0/0

Tags:

- Key: ProjectNumber

Value: 4

- Key: ProjectName

Value: interviewData

TestEIP:

Type: AWS::EC2::EIP

Properties:

Tags:

- Key: ProjectNumber

Value: 4

- Key: ProjectName

Value: interviewData

TestNetworkInterface:

Type: AWS::EC2::NetworkInterface

Properties:

Description: A External Network Interface for the EC2

SubnetId: !Ref TestPublicSubnet



GroupSet:

- !Ref TestInstanceSecurityGroup

Tags:

- Key: ProjectNumber

Value: 4

- Key: ProjectName

Value: interviewData

TestEIPAssociation:

Type: AWS::EC2::EIPAssociation

Properties:

AllocationId: !GetAtt TestEIP.AllocationId

NetworkInterfaceId: !Ref TestNetworkInterface

DependsOn: TestNetworkInterface

DependsOn: TestEIP

TestEC2Instance:

Type: 'AWS::EC2::Instance'

Properties:

**ImageId:** ami-002f6e91abff6eb96 # ami-053b12d3152c0cc71 for ubuntu

**InstanceType:** !Ref EC2InstanceSizeInput

**IamInstanceProfile** : !Ref InstanceProfileOfRoleToEC2

# SubnetId : !Ref TestPublicSubnet

# SecurityGroupIds:

# - !Ref TestInstanceSecurityGroup

**NetworkInterfaces:**

- **Description:** A Network interface made externally with AWS EIP attached at start up as primary

**DeviceIndex:** 0

**NetworkInterfaceId:** !Ref TestNetworkInterface

**KeyName** : myEC2KeyForInterview

**Tags:**

- **Key:** ProjectNumber

**Value:** 4

- **Key:** ProjectName

**Value:** interviewData

**UserData:**

**Fn::Base64:** !Sub |

#!/bin/bash

```
dnf update -y
```

```
dnf install httpd git python pip -y
```

```
yum install docker -y
```

```
systemctl start docker
```

```
systemctl enable docker
```

```
usermod -aG docker $USER
```

```
mkdir -p /usr/local/lib/docker/cli-plugins
```

```
curl -SL
```

```
https://github.com/docker/compose/releases/latest/download/docker-compose-linux-x86_64 -o /usr/local/lib/docker/cli-plugins/docker-compose
```

```
chmod +x /usr/local/lib/docker/cli-plugins/docker-compose
```

```
cd /home/ec2-user
```

```
git clone https://github.com/AryanGitHub/a-very-simple-webapp-for-assignment.git 2> error.log
```

```
cd a-very-simple-webapp-for-assignment
```

```
bash bash.sh 2> error_bash.log
```

**DependsOn:** TestEIPAssociation

**SSMEC2ControlRole:**

**Type:** AWS::IAM::Role

**Properties:**

**Description:** SSM Role for Test EC2

AssumeRolePolicyDocument:

Version: "2012-10-17"

Statement:

- Effect: Allow

Principal:

Service:

- ec2.amazonaws.com

Action:

- 'sts:AssumeRole'

ManagedPolicyArns:

- arn:aws:iam::aws:policy/AmazonSSMManagedInstanceCore

MaxSessionDuration: 3600

RoleName: Test\_EC2\_Role

Policies: # Adding inline policy for CloudWatch Logs

- PolicyName: CloudWatchLogsPolicy

PolicyDocument:

Version: "2012-10-17"

Statement:

- Effect: Allow

Action:

- logs:CreateLogGroup
- logs:CreateLogStream
- logs:PutLogEvents
- logs:DescribeLogStreams

Resource: "\*"

Tags:

- Key: ProjectNumber

Value: 4

- Key: ProjectName

Value: interviewData

InstanceProfileOfRoleToEC2:

Type: AWS::IAM::InstanceProfile

Properties:

InstanceProfileName: SSMEC2Role

Roles:

- !Ref SSMEC2ControlRole

### 3. Build the app and push it on github

main.py

```

from fastapi import FastAPI, Form, Request
from fastapi.responses import HTMLResponse, RedirectResponse
from fastapi.templating import Jinja2Templates
from prometheus_fastapi_instrumentator import Instrumentator

app = FastAPI()
templates = Jinja2Templates(directory="templates")


todos = []
Instrumentator().instrument(app).expose(app)
@app.get("/", response_class=HTMLResponse)
async def read_root(request: Request):
    return templates.TemplateResponse("index.html",
{"request": request, "todos": todos})

@app.post("/add", response_class=HTMLResponse)
async def add_todo(request: Request, task: str = Form(...)):
    todos.append(task)
    return RedirectResponse(url="/", status_code=303)

```

and templates/index.html

```

<!DOCTYPE html>
<html>
<head>
    <title>FastAPI ToDo App</title>
</head>
<body>
    <h1> ToDo List</h1>

    <form action="/add" method="post">
        <input type="text" name="task" placeholder="Enter a
task" required>
        <button type="submit">Add.</button>
    </form>

    <ul>
        {% for todo in todos %}
            <li>{{ todo }}</li>

```

```
        {% else %}
            <li>No tasks yet!</li>
        {% endfor %}
    </ul>
</body>
</html>
```

add a bash file to host it

```
#!/bin/bash
python -m venv .venv
source .venv/bin/activate
pip install -r ./requirements.txt
pip install prometheus-fastapi-instrumentator
uvicorn main:app --host 0.0.0.0 --port 80 --reload
```

### 3. Deploy Cloudformation template using aws cli command

and This themplate contains USERDATA, so it will automatically pull app from github repo

```
#!/bin/bash

aws cloudformation deploy --region ap-south-1 \

--template-file ./main.yaml \

--stack-name ec2forinterview \

--tags madeFromCLI=yes anotherTagForAllStackResources=okay \

--capabilities CAPABILITY_NAMED_IAM

# --no-execute-changeset
```

## 4.AWS configurations used (Resource Groups, Networking, etc

### CloudFormation Resources Summary

#### VPC and Networking Resources

##### 1. VPC

- **Logical ID:** TestVPC
- **Type:** AWS::EC2::VPC
- **Properties:**
  - CidrBlock: 10.0.0.0/16

##### 2. Internet Gateway

- **Logical ID:** TestIGW
- **Type:** AWS::EC2::InternetGateway

##### 3. VPC Gateway Attachment

- **Logical ID:** TestAttachGateway
- **Type:** AWS::EC2::VPCEGatewayAttachment
- **Properties:**
  - VpcId: !Ref TestVPC
  - InternetGatewayId: !Ref TestIGW

##### 4. Subnet

- **Logical ID:** TestPublicSubnet
- **Type:** AWS::EC2::Subnet
- **Properties:**
  - VpcId: !Ref TestVPC
  - AvailabilityZone: ap-south-1a
  - CidrBlock: 10.0.0.1/24

##### 5. Route Table

- **Logical ID:** TestRouteTable
- **Type:** AWS::EC2::RouteTable
- **Properties:**
  - VpcId: !Ref TestVPC



## 6. Route

- **Logical ID:** TestInternetPublicRoute
- **Type:** AWS::EC2::Route
- **Properties:**
  - RouteTableId: !Ref TestRouteTable
  - DestinationCidrBlock: 0.0.0.0/0
  - GatewayId: !Ref TestIGW

## 7. Subnet Route Table Association

- **Logical ID:** TestRouteTableToTestSubnetAssociation
- **Type:** AWS::EC2::SubnetRouteTableAssociation
- **Properties:**
  - RouteTableId: !Ref TestRouteTable
  - SubnetId: !Ref TestPublicSubnet

# Security Resources

## 8. Security Group

- **Logical ID:** TestInstanceSecurityGroup
- **Type:** AWS::EC2::SecurityGroup
- **Properties:**
  - VpcId: !Ref TestVPC
  - Ingress and Egress rules defined for SSH, HTTP, HTTPS, and all traffic.

# EC2 Resources

## 9. Elastic IP

- **Logical ID:** TestEIP
- **Type:** AWS::EC2::EIP

## 10. Network Interface

- **Logical ID:** TestNetworkInterface
- **Type:** AWS::EC2::NetworkInterface
- **Properties:**
  - SubnetId: !Ref TestPublicSubnet

- GroupSet: !Ref TestInstanceSecurityGroup

## 11. EIP Association

- **Logical ID:** TestEIPAssociation
- **Type:** AWS::EC2::EIPAssociation
- **Properties:**
  - AllocationId: !GetAtt TestEIP.AllocationId
  - NetworkInterfaceId: !Ref TestNetworkInterface

## 12. EC2 Instance

- **Logical ID:** TestEC2Instance
- **Type:** AWS::EC2::Instance
- **Properties:**
  - ImageId: ami-002f6e91abff6eb96
  - InstanceType: !Ref EC2InstanceSizeInput
  - NetworkInterfaces: !Ref TestNetworkInterface
  - UserData: Script for instance initialization.

## IAM Resources


### 13. IAM Role

- **Logical ID:** SSMEC2ControlRole
- **Type:** AWS::IAM::Role
- **Properties:**
  - AssumeRolePolicyDocument for EC2
  - ManagedPolicyArns:
    - arn:aws:iam::aws:policy/AmazonSSMManagedInstanceCore
  - Inline policy for CloudWatch Logs.

### 14. IAM Instance Profile

- **Logical ID:** InstanceProfileOfRoleToEC2
- **Type:** AWS::IAM::InstanceProfile
- **Properties:**
  - Roles: !Ref SSMEC2ControlRole


## Screenshots



## ToDo List

Enter a task

- No tasks yet!



## ToDo List

Enter a task

- hello
- I am Aryan

# Task2 CI/CD Pipeline Implementation

## CI/CD pipeline YAML using Github actions

```
name: remote ssh command

on: [push]

jobs:
```

```

build:
  name: Build
  runs-on: ubuntu-latest
  steps:
    - name: Executing remote SSH commands using pem file
      uses: appleboy/ssh-action@v1
      with:
        host: ${ secrets.HOST }
        username: ec2-user
        key : ${ secrets.EC2_SSH_KEY }
        port: 22
        script: |
          whoami
          echo "Deploying on EC2, logged IN"
          sudo chown -R ec2-user:ec2-user /home/ec2-user/a-
very-simple-webapp-for-assignment
          git config --global --add safe.directory
/home/ec2-user/a-very-simple-webapp-for-assignment
          cd /home/ec2-user/a-very-simple-webapp-for-
assignment
          git pull
          echo "Deployment script ran successfully!"

```

## Explanation of different pipeline stages

I have used single stage, to deploy the changes into the EC2 which is hosting the webapp

I used github action `appleboy/ssh-action@v1` its used to ssh into the ec2. it uses pem file contents into github repo secretes to deploy

I just run git pull from origin cus, `uvicorn --reload` is used to automatically reload the contents as the chags in the repo are detected.

## How environment variables/secrets are managed

Env Vars are managed using the github secrets for complete repo. there i have saved

the HOST public IP address for the ec2

the contents of pem file used to login into ec2 machine.

## Task3 Security & Compliance (ISO, GDPR, SOC 2)

### Security Risks Identification and Mitigation Strategies

#### 1. Insecure CI/CD Pipeline & Secrets Management

**Risk:** Exposure of sensitive information (API keys, tokens, database credentials, etc.) due to insecure storage of environment variables or incorrect configurations of the CI/CD pipeline.

**Mitigation Strategies:**

- **Secret Management:** Manage secrets securely using AWS Secrets Manager or AWS Systems Manager Parameter Store.
- **Encryption:** Ensure secrets stored in AWS services are encrypted in transit and at rest using AWS Key Management Service (KMS).
- **Isolation of Environments:** Limit exposure by using separate AWS accounts or Virtual Private Clouds (VPCs) for production, staging, and development environments.
- **Regular Auditing:** Use AWS CloudTrail and AWS Config to continuously monitor and audit CI/CD configurations.

**Compliance Alignment:**

- **ISO 27001:** Aligns with requirements for secure access controls and encryption policies.
- **SOC 2:** Addresses the security and confidentiality of information.
- **GDPR:** Mandates a high level of protection for personal data.

#### 2. Inadequate Access Control & Privilege Escalation

**Risk:** Overly permissive roles or policies can allow unauthorized access or privilege escalation, increasing the attack surface.

**Mitigation Strategies:**

- **Role-Based Access Control (RBAC):** Implement AWS IAM to enforce the principle of least privilege.

- **Multi-Factor Authentication (MFA):** Enforce MFA for users accessing AWS environments.
  - **Regular Reviews:** Regularly review IAM roles and permissions to ensure they meet evolving security requirements.
- Compliance Alignment:**
- **ISO 27001:** Requires strict access management and regular review of permissions.
  - **SOC 2:** Mandates controls to protect operational environments from unauthorized access.
  - **GDPR:** Least-privilege access helps mitigate the risk of unauthorized personal data exposure.

### 3. Third-Party & Supply Chain Vulnerabilities

**Risk:** Using untrusted third-party libraries, container images, or external plugins can introduce vulnerabilities.

**Mitigation Strategies:**

- **Vulnerability Scanning:** Regularly scan dependencies and container images using tools like AWS Inspector or Amazon ECR.
- **Trusted Registries and Code Signing:** Use trusted registries for container images and implement code signing.
- **Update Policies:** Maintain an effective patch and update

## Task 4 Monitoring & Logging

### Steps to configure monitoring/logging tools

I have used Prometheus & Grafana

To set it up on the webapp used for task 1 and task 2, I made USERDATA of EC@ to install docker

add the configuration file for Prometheus

```
global:
  scrape_interval: 4s

scrape_configs:
  - job_name: prometheus
```

```
static_configs:
  - targets: ["10.0.0.36:80"]
```

and docker-compose file to run it with the given configurations

```
version: "3"

services:
  prom-server:
    image: prom/prometheus
    ports:
      - 9090:9090
    volumes:
      - ./prometheus-
config.yml:/etc/prometheus/prometheus.yml
```

then docker compose up and check

```
curl http://10.0.0.36:80/metrics
```

user http://:9090 to open Prometheus

For Grafana

run docker container

```
docker run -d -p 3000:3000 --name=grafana grafana/grafana-oss
```

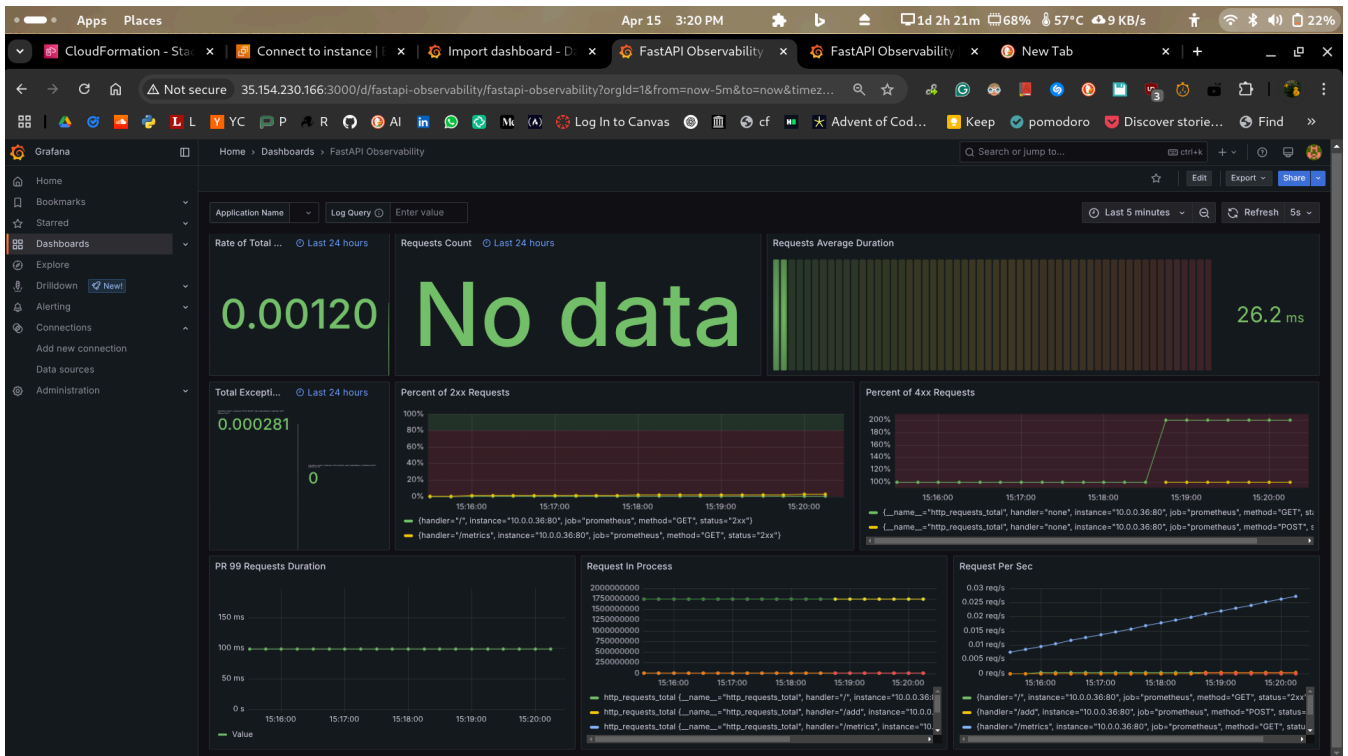
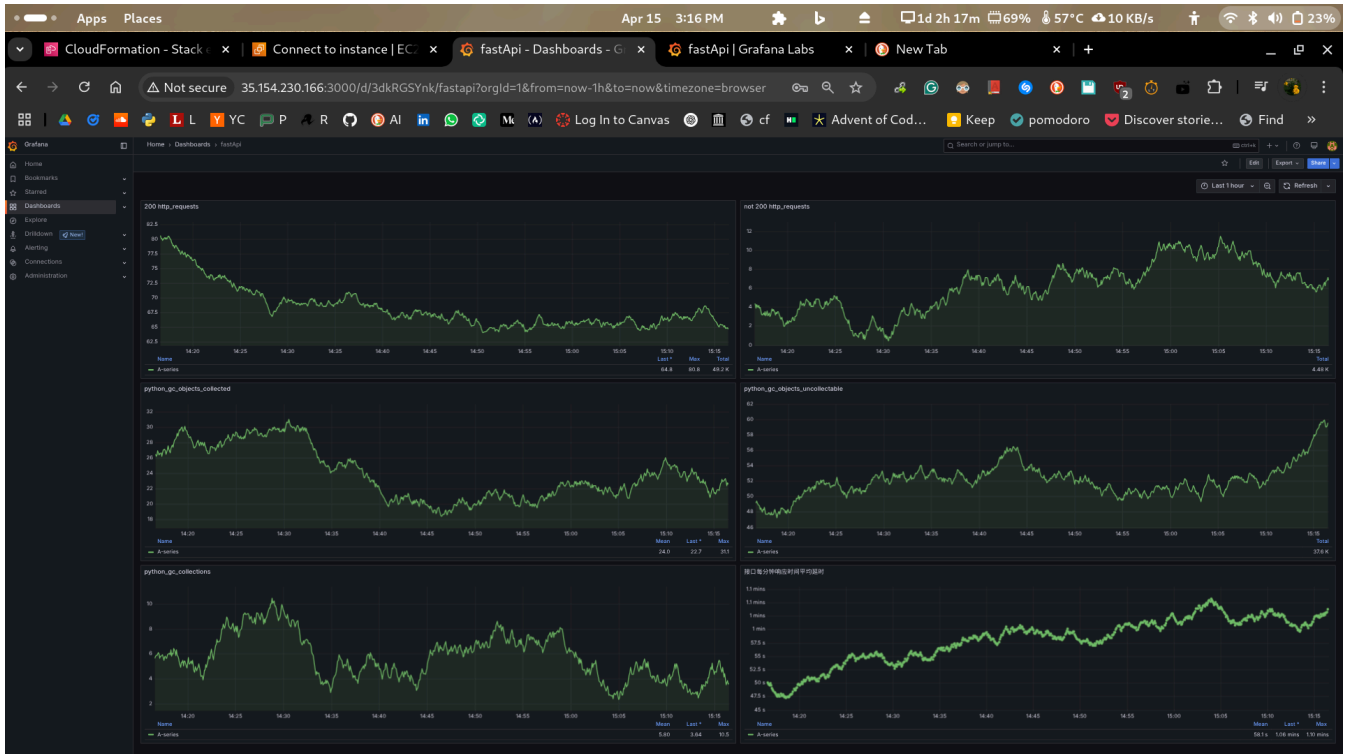
add data source to Prometheus  
by adding

```
http://10.0.0.36:9090
```


then login to grafana , and goto Dashboard > New dashboard > add id  
15834

we can add another dashboard, add id 18739

# Dashboard screenshots showing application metrics







# Task5 Database & Storage Optimization

## Overview

Database optimization enhances query performance, resource utilization, and overall reliability. These techniques are applicable to AWS-managed databases like Amazon RDS or Amazon DocumentDB.

## Optimization Techniques

### 1. Indexing

**Purpose:** Improve data retrieval times by reducing the amount of data scanned during query execution.

**Implementation:** Identify columns frequently used in query `WHERE` or `ORDER BY` clauses and create indexes on them.

**Example (PostgreSQL):**

**Before Indexing:**

```
SELECT * FROM orders WHERE customer_id = '12345' ORDER BY order_date DESC;
```

**After Indexing:**

```
CREATE INDEX idx_orders_customer_date ON orders
(customer_id, order_date DESC);
```

**Benefit:** This composite index allows direct access to relevant rows, improving query efficiency.

## 2. Query Optimization

**Purpose:** Rewrite queries to eliminate redundant computations and optimize join logic.

**Implementation:** Use tools like PostgreSQL's `EXPLAIN ANALYZE` to identify performance bottlenecks.

**Example:**

**Before (Using Subquery):**

```
SELECT * FROM orders
WHERE customer_id IN (SELECT id FROM customers WHERE
region = 'West');
```

**After (Using JOIN):**

```
SELECT o.*
FROM orders o
INNER JOIN customers c ON o.customer_id = c.id
WHERE c.region = 'West';
```

**Benefit:** Better performance due to more efficient join algorithm selection.

## 3. Data Partitioning

**Purpose:** Divide large tables into smaller, more manageable pieces to improve performance.

**Implementation:** Partition data based on a specific key, such as date ranges or categorical values.

**Example (PostgreSQL - Range Partitioning):**

```
-- Define the parent partitioned table
CREATE TABLE orders (
  order_id serial NOT NULL,
```

```

customer_id int NOT NULL,
order_date date NOT NULL,
-- other columns
PRIMARY KEY (order_id, order_date)
) PARTITION BY RANGE (order_date);

-- Create partitions
CREATE TABLE orders_2024 PARTITION OF orders
FOR VALUES FROM ('2024-01-01') TO ('2025-01-01');

CREATE TABLE orders_2025 PARTITION OF orders
FOR VALUES FROM ('2025-01-01') TO ('2026-01-01');

```

**Benefit:** Queries can scan only relevant partitions, significantly improving performance.

## Task 6 Automation & Scripting

### Script file (.sh or .py)

```

#!/bin/bash
# deploy the cloudformation file
aws cloudformation deploy --region ap-south-1 \

--template-file ./main.yaml \

--stack-name ec2forinterview \

--tags madeFromCLI=yeah anotherTagForAllStackResources=okay \

--capabilities CAPABILITY_NAMED_IAM

# --no-execute-changeset

```

### Explanation of what the script does

This simple commands build and deploy the cloudformation template, the template contains the user data

```
#!/bin/bash
```

```
dnf update -y
```

```
dnf install httpd git python pip -y
```

```
yum install docker -y
```

```
systemctl start docker
```

```
systemctl enable docker
```

```
usermod -aG docker $USER
```

```
mkdir -p /usr/local/lib/docker/cli-plugins
```

```
curl -SL
```

```
https://github.com/docker/compose/releases/latest/download/docker-compose-linux-x86_64 -o /usr/local/lib/docker/cli-plugins/docker-compose
```

```
chmod +x /usr/local/lib/docker/cli-plugins/docker-compose
```

```
cd /home/ec2-user
```

```
git clone https://github.com/AryanGitHub/a-very-simple-webapp-for-assignment.git 2> error.log
```

```
cd a-very-simple-webapp-for-assignment
```

```
bash bash.sh 2> error_bash.log
```

bash file in the webapp folder

```
#!/bin/bash
```

```
python -m venv .venv
```

```
source .venv/bin/activate
```

```
pip install -r ./requirements.txt
```

```
pip install prometheus-fastapi-instrumentator
uvicorn main:app --host 0.0.0.0 --port 80 --reload
```

it installs the python , git , pip, docker and docker compose  
then it clones the repo  
and starts running the bash script which runs the webapp, the script  
download all the packages from requiremet.txt file in an virtual environment  
and then runs the uvicorn command

## Task7 Disaster Recovery & High Availability

### Key Elements of a DR Strategy

#### 1. Recovery Time Objective (RTO)

**Definition:** Maximum acceptable downtime duration before service restoration.

**Example Target:** 2 hours (system must be fully functional within 2 hours of incident).

#### 2. Recovery Point Objective (RPO)

**Definition:** Maximum acceptable data loss measured in time before failure.

**Example Target:** 15 minutes (data loss should not exceed the last 15 minutes of transactions).

#### 3. Backup Strategy

- **Automated Backups:** Schedule regular backups using AWS Backup or native service features.
- **Geographical Redundancy:** Store critical backups in a separate AWS Region.
- **Testing and Validation:** Regularly test restore procedures to validate RTO and RPO targets.
- **Backup Security:** Encrypt all backups using AWS KMS and restrict access using IAM policies.

#### 4. High Availability (HA) Implementation

- **Multi-AZ Deployments:** Deploy across multiple Availability Zones within an AWS Region.

- **Replication:** Utilize synchronous or near-synchronous replication for critical data.
- **Load Balancing:** Use Elastic Load Balancing to distribute traffic across healthy instances.
- **Automated Failover:** Configure services like Amazon RDS Multi-AZ for automatic failover.

## Example: Automated Backup Setup using AWS Backup

### 1. Create a Backup Vault:

```
aws backup create-backup-vault \  
  --backup-vault-name MyBackupVault \  
  --region us-east-1
```

### 2. Define and Assign a Backup Plan:

Create a JSON file (backup-plan.json):

```
{  
  "BackupPlanName": "DailyBackupPlan",  
  "Rules": [  
    {  
      "RuleName": "DailyFullBackup",  
      "TargetBackupVaultName": "MyBackupVault",  
      "ScheduleExpression": "cron(0 2 * * ? *)",  
      "StartWindowMinutes": 60,  
      "CompletionWindowMinutes": 180,  
      "Lifecycle": {  
        "DeleteAfterDays": 30  
      }  
    }  
  ]  
}
```

Create the backup plan:

```
aws backup create-backup-plan --backup-plan file://backup-plan.json
```

Assign resources:

```
aws backup create-backup-selection \
  --backup-plan-id <your-backup-plan-id> \
  --backup-selection '{
    "SelectionName": "EC2Selection",
    "IamRoleArn": "arn:aws:iam::<account-id>:role/service-role/AWSBackupDefaultServiceRole",
    "Resources": [
      "arn:aws:ec2:us-east-1:<account-id>:instance/<instance-id>"
    ]
  }' \
  --creator-request-id $(uuidgen)
```

## Example: Setting Up Automated Backups in AWS

AWS offers automated backup solutions integrated with many of its services. For instance, using AWS Backup you can centralize the backup of EC2 instances, RDS databases, and more. Below is an example using AWS CLI commands:

### 1. Create a Backup Vault:

```
aws backup create-backup-vault \
  --backup-vault-name MyBackupVault \
  --region us-east-1
```

This command creates a backup vault in the specified region to store your backup data.

### 2. Create and Assign a Backup Plan:

First, create a JSON file (e.g., backup-plan.json) with your backup plan details:

```
{
  "BackupPlanName": "DailyBackupPlan",
  "Rules": [
    {
      "RuleName": "DailyFullBackup",
      "TargetBackupVaultName": "MyBackupVault",
      "ScheduleExpression": "cron(0 2 * * ? *)",
      "StartWindowMinutes": 60,
      "CompletionWindowMinutes": 180,
      "Lifecycle": {
        "DeleteAfterDays": 30
      }
    }
  ]
}
```

Then, apply the backup plan and assign resources (for example, an EC2 instance):

```
# Create the backup plan
aws backup create-backup-plan --backup-plan file://backup-plan.json

# Assign resources (example for EC2, ensuring the resource ARN is correct)
aws backup create-backup-selection \
  --backup-plan-id <your-backup-plan-id> \
  --backup-selection '{
    "SelectionName": "EC2Selection",
    "IamRoleArn": "arn:aws:iam::<account-id>:role/AWSBackupDefaultServiceRole",
    "Resources": ["arn:aws:ec2:us-east-1:<account-id>:instance/<instance-id>"]
  }'
```

This setup schedules a daily full backup for the

## Task8 AI Model Deployment & MLOps



We have used ECS to host the docker task, and added ALB (Application load balancer to the publically exposed docker service)

This all is achieved using cloudformation template and the ECS Fargate.

The docker service takes the voice sound file and convet it into test.

The famous docker image used for this is onerahmet/openai-whisper-asr-webbservice:latest

## Dockerfile & Kubernetes YAML files

### Cloudformation File

```
AWSTemplateFormatVersion: '2010-09-09'

Description: Deploy Whisper ASR API to ECS Fargate
Parameters:

  WhisperModel:

    Type: String

    Default: tiny

    AllowedValues: [tiny, base, small, medium, large]

    Description: Whisper model to use

Resources:

  WhisperVPC:

    Type: AWS::EC2::VPC

    Properties:
```

CidrBlock: 10.0.0.0/16

EnableDnsSupport: **true**

EnableDnsHostnames: **true**

Tags: [{ Key: Name, Value: WhisperVPC }]

WhisperSubnet1:

Type: AWS::EC2::Subnet

Properties:

VpcId: **!Ref** WhisperVPC

CidrBlock: 10.0.1.0/24

AvailabilityZone: **!Select** [0, **!GetAZs** '']

MapPublicIpOnLaunch: **true**

WhisperSubnet2:

Type: AWS::EC2::Subnet

Properties:

VpcId: **!Ref** WhisperVPC

CidrBlock: 10.0.2.0/24

AvailabilityZone: **!Select** [1, **!GetAZs** '']

MapPublicIpOnLaunch: **true**

WhisperInternetGateway:

Type: AWS::EC2::InternetGateway

WhisperAttachGateway:

Type: AWS::EC2::VPCGatewayAttachment

Properties:

VpcId: !Ref WhisperVPC

InternetGatewayId: !Ref WhisperInternetGateway

WhisperRouteTable:

Type: AWS::EC2::RouteTable

Properties:

VpcId: !Ref WhisperVPC

WhisperRoute:

Type: AWS::EC2::Route

DependsOn: WhisperAttachGateway

Properties:

RouteTableId: !Ref WhisperRouteTable

**DestinationCidrBlock:** 0.0.0.0/0

**GatewayId:** !Ref WhisperInternetGateway

**WhisperSubnetRouteTableAssoc1:**

**Type:** AWS::EC2::SubnetRouteTableAssociation

**Properties:**

**SubnetId:** !Ref WhisperSubnet1

**RouteTableId:** !Ref WhisperRouteTable

**WhisperSubnetRouteTableAssoc2:**

**Type:** AWS::EC2::SubnetRouteTableAssociation

**Properties:**

**SubnetId:** !Ref WhisperSubnet2

**RouteTableId:** !Ref WhisperRouteTable

**WhisperSecurityGroup:**

**Type:** AWS::EC2::SecurityGroup

**Properties:**

**GroupDescription:** Allow HTTP access

**VpcId:** !Ref WhisperVPC

SecurityGroupIngress:

- IpProtocol: tcp

FromPort: 9000

ToPort: 9000

CidrIp: 0.0.0.0/0

WhisperCluster:

Type: AWS::ECS::Cluster

WhisperTaskExecutionRole:

Type: AWS::IAM::Role

Properties:

AssumeRolePolicyDocument:

Statement:

- Effect: Allow

Principal:

Service: ecs-tasks.amazonaws.com

Action: sts:AssumeRole

ManagedPolicyArns:

- arn:aws:iam::aws:policy/service-

role/AmazonECSTaskExecutionRolePolicy

WhisperTaskDefinition:

Type: AWS::ECS::TaskDefinition

Properties:

Family: whisper-task

RequiresCompatibilities: [FARGATE]

Cpu: 512

Memory: 1024

NetworkMode: awsvpc

ExecutionRoleArn: !GetAtt WhisperTaskExecutionRole.Arn

ContainerDefinitions:

- Name: whisper

Image: onerahmet/openai-whisper-asr-webservice:latest

PortMappings:

- ContainerPort: 9000

Environment:

- Name: ASR\_MODEL

Value: !Ref WhisperModel

WhisperService:

Type: AWS::ECS::Service

DependsOn: WhisperALBListener

Properties:

Cluster: !Ref WhisperCluster

LaunchType: FARGATE

DesiredCount: 1

NetworkConfiguration:

AwsvpcConfiguration:

AssignPublicIp: ENABLED

SecurityGroups: [!Ref WhisperSecurityGroup]

Subnets: [!Ref WhisperSubnet1, !Ref WhisperSubnet2]

TaskDefinition: !Ref WhisperTaskDefinition

LoadBalancers:

- ContainerName: whisper

ContainerPort: 9000

TargetGroupArn: !Ref WhisperTargetGroup

WhisperALB:

Type: AWS::ElasticLoadBalancingV2::LoadBalancer

### Properties:

Name: whisper-alb

Subnets: [!Ref WhisperSubnet1, !Ref WhisperSubnet2]

SecurityGroups: [!Ref WhisperSecurityGroup]

Scheme: internet-facing

Type: application

### WhisperTargetGroup:

Type: AWS::ElasticLoadBalancingV2::TargetGroup

### Properties:

Port: 9000

Protocol: HTTP

VpcId: !Ref WhisperVPC

TargetType: ip

HealthCheckPath: /docs

### WhisperALBListener:

Type: AWS::ElasticLoadBalancingV2::Listener

### Properties:

LoadBalancerArn: !Ref WhisperALB



Port: 9000

Protocol: HTTP

DefaultActions:

- Type: forward

TargetGroupArn: !Ref WhisperTargetGroup

Outputs:

WhisperAPIURL:

Description: Whisper REST API URL

Value: !Join ["", ["http://", !GetAtt WhisperALB.DNSName, ":9000"]]

## Steps to deploy the model

deploy this using AWS CLI deploy command

```
#!/bin/bash
```

```
aws cloudformation deploy --region ap-south-1 \
```

```
--template-file ./main.yaml \
```

```
--stack-name ecsaimodel \
```

```
--tags madeFromCLI=yeah anotherTagForAllStackResources=okay \
```

```
--capabilities CAPABILITY_NAMED_IAM
```

```
# --no-execute-changeset
```

# Screenshot of the model running on ECS

The screenshot shows the AWS Management Console for the 'Whisper Asr Web Service'. The 'Endpoints' tab is selected, showing the 'asr' endpoint. The endpoint is a POST request to 'http://whisper-alb-765361359.ap-south-1.elb.amazonaws.com:9000/docs#/Endpoints/asr\_asr\_post'. The endpoint parameters are:

- encode** (boolean, query): Encode audio first through ffmpeg. Default value: true. Value: true.
- task** (string, query): Available values: transcribe, translate. Default value: transcribe. Value: transcribe.
- language** (string, query): Available values: af, am, ar, as, az, ba, be, bg, bn, bo, br, bs, ca, cs, la, lb, ln, lo, lt, lv, mg, mi, mk, ml, mn, mr, ms, mt, my, ne, nl, nn, no, oc, pa, pl, ps, pt, ro, ru, sa, sd, si, sk, sl, sn, so, sq, sr, su, sv, sw, ta, te, tg, th, tk, tl, tr, tt, uk, ur, uz, vi, yi, yo, yue, zh. Value: --.
- initial\_prompt** (string, query): initial\_prompt. Value: initial\_prompt.
- output** (string, query): Available values: txt, vtt, srt, tsv, json. Default value: txt. Value: txt.

The 'Request body' section is marked as 'required' and shows 'audio\_file \* required' and 'string(\$binary)'.

A terminal window in the foreground shows the following command and output:

```
ooumua@fedora: ~/Downloads/test1
[ooumua@fedora test1]$ curl -X POST http://whisper-alb-765361359.ap-south-1.elb.amazonaws.com:9000/asr \
-F audio_file=@rec1.flac
Hello world, this is Aryan Pandey. Thank you.
[ooumua@fedora test1]$
```

The screenshot shows the OpenAPI specification for the 'Whisper Asr Web Service'. The 'Endpoints' section is expanded, showing the 'asr' endpoint. The endpoint is a POST request to 'http://whisper-alb-765361359.ap-south-1.elb.amazonaws.com:9000/docs#/Endpoints/asr\_asr\_post'. The endpoint parameters are:

- encode** (boolean, query): Encode audio first through ffmpeg. Default value: true. Value: true.
- task** (string, query): Available values: transcribe, translate. Default value: transcribe. Value: transcribe.
- language** (string, query): Available values: af, am, ar, as, az, ba, be, bg, bn, bo, br, bs, ca, cs, cy, da, de, el, en, es, et, eu, fa, fi, fo, fr, gl, gu, ha, haw, he, hi, hr, ht, hu, hy, id, is, it, ja, jw, ka, kk, km, kn, ko, la, lb, ln, lo, lt, lv, mg, mi, mk, ml, mn, mr, ms, mt, my, ne, nl, nn, no, oc, pa, pl, ps, pt, ro, ru, sa, sd, si, sk, sl, sn, so, sq, sr, su, sv, sw, ta, te, tg, th, tk, tl, tr, tt, uk, ur, uz, vi, yi, yo, yue, zh. Value: --.
- initial\_prompt** (string, query): initial\_prompt. Value: initial\_prompt.

The 'Request body' section is marked as 'required' and shows 'audio\_file \* required' and 'string(\$binary)'.

A terminal window in the foreground shows the following command and output:

```
ooumua@fedora: ~/Downloads/test1
[ooumua@fedora test1]$ curl -X POST http://whisper-alb-765361359.ap-south-1.elb.amazonaws.com:9000/asr \
-F audio_file=@rec1.flac
Hello world, this is Aryan Pandey. Thank you.
[ooumua@fedora test1]$
```

Amazon Elastic Container Service

Clusters

Namespaces

Task definitions

Account settings

Install AWS Copilot

Amazon ECR

Repositories

AWS Batch

Documentation

Discover products

Subscriptions

Service overview

Status

Active

Tasks (1 Desired)

0 Pending | 1 Running

Task definition: revision

whisper-task:1

Deployment status

Success

Health and metrics

Tasks

Logs

Deployments

Events

Configuration and networking

Service auto scaling

Tags

Status

Service name

ecsaimodel-WhisperService-nM1UjeswHwLW

Service ARN

arn:aws:ecs:ap-south-1:222634371739:service/ecsaimodel-WhisperCluster-6QbnxBNgKJXh/ecsaimodel-WhisperService-nM1UjeswHwLW

Deployments current state

1 Completed task

Created at

April 15, 2025 at 02:52 (UTC+5:30)

Health check grace period

0 seconds

Load balancer health

Load balancer

Load balancer type

Listeners

Target group

Targets

whisper-alb

Application Load Balancer

HTTP:9000

ecsaim-Whisp-UFGBACU1I3XX

Details

1 Healthy | 0 Unhealthy

CloudShell

Feedback

© 2025, Amazon Web Services, Inc. or its affiliates.

Privacy

Terms

Cookie preferences

EC2

Dashboard

EC2 Global View

Events

Instances

Instance Types

Launch Templates

Spot Requests

Savings Plans

Reserved Instances

Dedicated Hosts

Capacity Reservations

Images

AMIs

AMI Catalog

Elastic Block Store

Volumes

Snapshots

whisper-alb

Details

Load balancer type

Application

Scheme

Internet-facing

Status

Active

Hosted zone

Zone ID: ZP97RAFLXTNZK

VPC

vpc-039de96d703ff2ba8

Availability Zones

subnet-07c5088d0eba5ecee ap-south-1a (aps1-az1)  
subnet-0753968c972545143 ap-south-1b (aps1-az3)

Load balancer IP address type

IPv4

Date created

April 15, 2025, 02:49 (UTC+05:30)

Load balancer ARN

arn:aws:elasticloadbalancing:ap-south-1:222634371739:loadbalancer/app/whisper-alb/ac8abdccb0a6e001

DNS name

whisper-alb-765361359.ap-south-1.elb.amazonaws.com (A Record)

Listeners and rules (1)

Network mapping

Resource map

Security

Monitoring

Integrations

Attributes

Capacity

Tags

Listeners and rules (1)

A listener checks for connection requests on its configured protocol and port. Traffic received by the listener is routed according to the default action and any additional rules.

Filter listeners

1

CloudShell

Feedback

© 2025, Amazon Web Services, Inc. or its affiliates.

Privacy

Terms

Cookie preferences

CloudFormation console showing the 'Stacks' view for the 'ecsaiproject' stack. The stack is in the 'Active' state and shows a list of resources (17) that are all in the 'CREATE\_COMPLETE' state.

**Stacks (5)**

- ecsaiproject (2025-04-15 02:48:41 UTC+0530) **CREATE\_COMPLETE**
- filescanRestApi (2025-04-10 22:00:53 UTC+0530) **CREATE\_COMPLETE**
- filescanBackend (2025-04-10 21:58:16 UTC+0530) **CREATE\_COMPLETE**
- malwareGuardianFrontendStack (2025-04-10 21:52:07 UTC+0530) **CREATE\_COMPLETE**

**Resources (17)**

Logical ID	Physical ID	Type	Status
WhisperALB	arn:aws:elasticloadbalancing:ap-south-1:222634371739:loadbalancer/app/whisper-alb/ac8abdcdb0a6e001	AWS::ElasticLoadBalancingV2::LoadBalancer	CREATE_COMPLETE
WhisperALBListener	arn:aws:elasticloadbalancing:ap-south-1:222634371739:listener/app/whisper-alb/ac8abdcdb0a6e001/67f2e7c236028d33	AWS::ElasticLoadBalancingV2::Listener	CREATE_COMPLETE
WhisperAttachGateway	IGWVjpc-039de96d703ff2ba8	AWS::EC2::VPCGatewayAttachment	CREATE_COMPLETE
WhisperCluster	ecsaiproject-WhisperCluster-6QbnxBNgKJXh	AWS::ECS::Cluster	CREATE_COMPLETE

CloudFormation console showing the 'Stacks' view for the 'ecsaiproject' stack. The stack is in the 'Active' state and shows a list of resources (17) that are all in the 'CREATE\_COMPLETE' state.

**Stacks (5)**

- ecsaiproject (2025-04-15 02:48:41 UTC+0530) **CREATE\_COMPLETE**
- filescanRestApi (2025-04-10 22:00:53 UTC+0530) **CREATE\_COMPLETE**
- filescanBackend (2025-04-10 21:58:16 UTC+0530) **CREATE\_COMPLETE**
- malwareGuardianFrontendStack (2025-04-10 21:52:07 UTC+0530) **CREATE\_COMPLETE**

**Resources (17)**

Logical ID	Physical ID	Type	Status
WhisperService		AWS::ECS::Service	CREATE_COMPLETE
WhisperALBListener		AWS::ElasticLoadBalancingV2::Listener	CREATE_COMPLETE
WhisperRoute		AWS::ElasticLoadBalancingV2::Rule	CREATE_COMPLETE
WhisperALB		AWS::ElasticLoadBalancingV2::LoadBalancer	CREATE_COMPLETE
WhisperTaskDefinition		AWS::ECS::TaskDefinition	CREATE_COMPLETE
WhisperSubnetRouteTableAs...		AWS::EC2::RouteTable	CREATE_COMPLETE
WhisperSubnetRoute TableAs...		AWS::EC2::RouteTable	CREATE_COMPLETE
WhisperAttachGateway		AWS::EC2::VPCGatewayAttachment	CREATE_COMPLETE
WhisperSecurityGroup		AWS::EC2::SecurityGroup	CREATE_COMPLETE
WhisperSubnet1		AWS::EC2::Subnet	CREATE_COMPLETE
WhisperTargetGroup		AWS::ElasticLoadBalancingV2::TargetGroup	CREATE_COMPLETE
WhisperSubnet2		AWS::EC2::Subnet	CREATE_COMPLETE
WhisperRoute Table		AWS::EC2::RouteTable	CREATE_COMPLETE
WhisperVPC		AWS::EC2::VPC	CREATE_COMPLETE
WhisperInternetGateway		AWS::EC2::InternetGateway	CREATE_COMPLETE
WhisperCluster		AWS::ECS::Cluster	CREATE_COMPLETE
WhisperTaskExecutionRole		AWS::IAM::Role	CREATE_COMPLETE